

Code

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report
```

```
from sklearn.preprocessing import LabelEncoder
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
#load the dataset
```

```
data=pd.read_csv('titanic_sample.csv')
```

```
print(data)
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35	0	0	373450	8.0500	NaN	S

```
#explore the dataset
```

```
print("Dataset Head:\n",data.head())
```

Dataset Head:												
	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282	7.9250	NaN	S
3				Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803	53.1000	C123	S
4				Allen, Mr. William Henry	male	35	0	0	373450	8.0500	NaN	S

```
#explore the dataset
```

```
print("\nDataset Info",data.info())
```

```
RangeIndex: 5 entries, 0 to 4
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   PassengerId  5 non-null      int64
 1   Survived     5 non-null      int64
 2   Pclass       5 non-null      int64
 3   Name         5 non-null      object
 4   Sex          5 non-null      object
 5   Age          5 non-null      int64
 6   SibSp        5 non-null      int64
 7   Parch        5 non-null      int64
 8   Ticket       5 non-null      object
 9   Fare         5 non-null      float64
10   Cabin        2 non-null      object
11   Embarked     5 non-null      object
dtypes: float64(1), int64(6), object(5)
memory usage: 612.0+ bytes
```

```
print("\nMissing Values:\n",data.isnull().sum())
```


```
Missing Values:
 PassengerId    0
 Survived       0
 Pclass         0
 Name           0
 Sex            0
 Age            0
 SibSp          0
 Parch          0
 Ticket         0
 Fare           0
 Cabin          3
 Embarked       0
dtype: int64
```

```
#Visualize some features
```

```
sns.countplot(x='Survived',data=data)
```

```
plt.title('Survival Count')
```

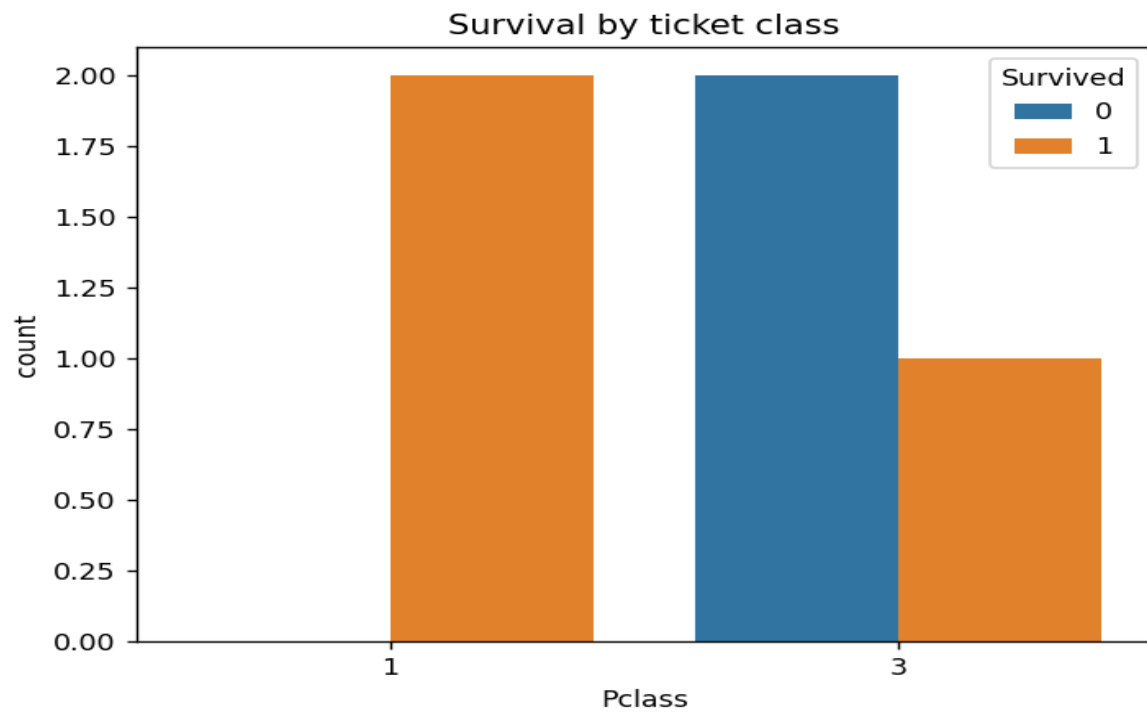
```
plt.show()
```



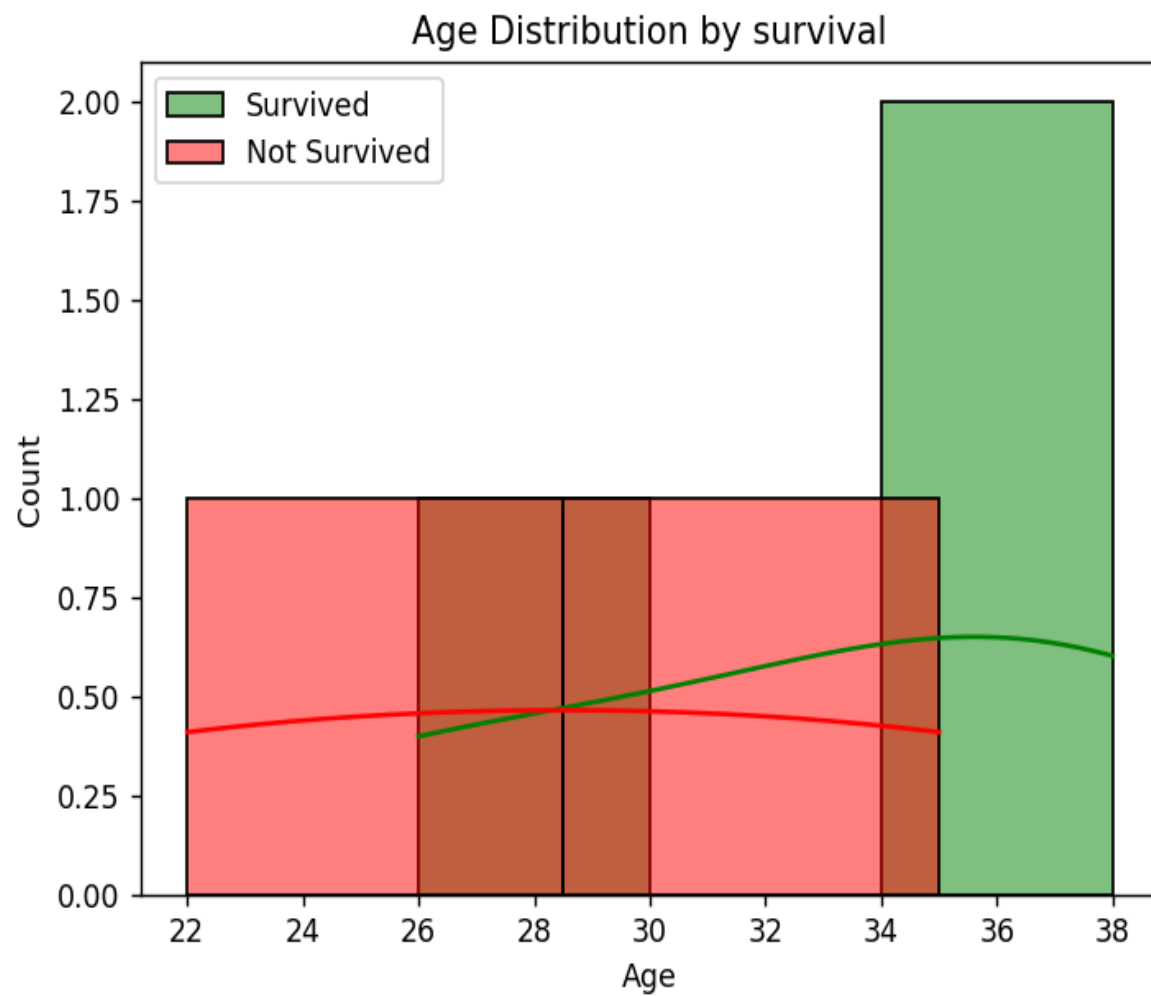
```
sns.countplot(x='Pclass',hue='Survived',data=data)
```

```
plt.title('Survival by ticket class')
```

```
plt.sho()
```



```
sns.histplot(data[data['Survived']==1]['Age'],kde=True,label='Survived',color='green')
sns.histplot(data[data['Survived']==0]['Age'],kde=True,label='Not Survived',color='red')
plt.legend()
plt.title("Age Distribution by survival")
plt.show()
```



```
#data cleaning
```

```
#fill missing Age values with the median
```

```
data['Age'].fillna(data['Age'].median(),inplace=True)
```

```
print(data['Age'])
```

```
data['Age'].fillna(data['Age'].median(),inplace=True)
0    22
1    38
2    26
3    35
4    35
Name: Age, dtype: int64
```

```
#filling missing Embarked values with the mode
```

```
data['Embarked'].fillna(data['Embarked'].mode()[0],inplace=True)
```

```
print(data['Embarked'])
```

```
data['Embarked'].fillna(data['Embarked'].mode()[0],inplace=True)
0    S
1    C
2    S
3    S
4    S
Name: Embarked, dtype: object
```

```
#drop cabin(too many missin values)
```

```
data.drop('Cabin',axis=1,inplace=True)
```

```
print(data.drop)
```

```
<bound method DataFrame.drop of
PassengerId  Survived  Pclass
0            1         0       3  Braund, Mr. Owen Harris    male  22    1    0    A/5 21171  7.2500  S
1            2         1       1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38    1    0    PC 17599  71.2833  C
2            3         1       3                Heikinen, Miss. Laina female  26    0    0  STON/O2. 3101282  7.9250  S
3            4         1       1  Futrelle, Mrs. Jacques Heath (Lily May Peel) female  35    1    0          113803  53.1000  S
4            5         0       3                Allen, Mr. William Henry    male  35    0    0          373450  8.0500  S>
```

```
#Drop irrelevant features
```

```
data.drop(['Name','Ticket','PassengerId'],axis=1,inplace=True)
```

```
print(data.drop)
```

```
<bound method DataFrame.drop of
0      0      3  male  22      1      0  7.2500      S
1      1      1  female 38      1      0 71.2833      C
2      1      3  female 26      0      0  7.9250      S
3      1      1  female 35      1      0 53.1000      S
4      0      3  male  35      0      0  8.0500      S>
```

```
#Encode categorical variables
```

```
encoder=LabelEncoder()
```

```
data['Sex']=encoder.fit_transform(data['Sex'])
```

```
data['Embarked']=encoder.fit_transform(data['Embarked'])
```

```
print(data['Sex'])
```

```
print(data['Embarked'])
```

```
0      1
1      0
2      0
3      0
4      1
Name: Sex, dtype: int64
0      1
1      0
2      1
3      1
4      1
Name: Embarked, dtype: int64
```

```
#Feature Engineering
```

```
#Create familySize feature
```

```
data['FamilySize']=data['SibSp'] + data['Parch'] +1
```

```
print(data['FamilySize'])
```

```
data.drop(['SibSp','Parch'],axis=1,inplace=True)
```

```
print(data.drop)
```

```
<bound method DataFrame.drop of      Survived  Pclass  Sex  Age      Fare  Embarked  FamilySize
0           0       3    1  22   7.2500         1         2
1           1       1    0  38  71.2833         0         2
2           1       3    0  26   7.9250         1         1
3           1       1    0  35  53.1000         1         2
4           1       3    1  35   8.0500         1         1>
```

```
#split the data into features and target
```

```
x=data.drop('Survived',axis=1)
```

```
print(x)
```

```
y=data['Survived']
```

```
print(y)
```

```
      Pclass  Sex  Age      Fare  Embarked  FamilySize
0         3    1  22   7.2500         1         2
1         1    0  38  71.2833         0         2
2         3    0  26   7.9250         1         1
3         1    0  35  53.1000         1         2
4         3    1  35   8.0500         1         1
0         0
1         1
2         1
3         1
4         0
Name: Survived, dtype: int64
```



```
#train-test split
```

```
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
print(x_train,x_test,y_train,y_test)
```

```
   Pclass  Sex  Age   Fare  Embarked  FamilySize
4        3    1   35   8.050         1           1
2        3    0   26   7.925         1           1
0        3    1   22   7.250         1           2
3        1    0   35  53.100         1           2
1        1    0   38  71.2833         0           2
2        1
0        0
3        1
Name: Survived, dtype: int64 1    1
Name: Survived, dtype: int64
```

```
#Build and train models
```

```
#Logistic Regression
```

```
log_reg=LogisticRegression(max_iter=1000)
```

```
print(log_reg)
```

```
log_reg.fit(x_train,y_train)
```

```
print(log_reg.fit)
```

```
LogisticRegression(max_iter=1000)
<bound method LogisticRegression.fit of LogisticRegression(max_iter=1000)>
```

```
#Random Forest
```

```
rf=RandomForestClassifier(random_state=42)
```

```
rf.fit(x_train,y_train)
```

```
print(rf)
```

```
print(rf.fit)
```

```
RandomForestClassifier(random_state=42)
<bound method BaseForest.fit of RandomForestClassifier(random_state=42)>
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
classification_report
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.linear_model import LogisticRegression
```

```
# Assume x_train, x_test, y_train, y_test are already defined
```

```
# Define models
```

```
log_reg = LogisticRegression(random_state=42)
```

```
print(log_reg)
```

```
rf = RandomForestClassifier(random_state=42)
```

```
print(rf)
```

```
LogisticRegression(random_state=42)
RandomForestClassifier(random_state=42)
```

```
# Train models
```

```
log_reg.fit(x_train, y_train)
```

```
rf.fit(x_train, y_train)
```

```
# Evaluate models
```

```
models = {'Logistic Regression': log_reg, 'Random Forest': rf}
```

```
for name, model in models.items():
```

```
    y_pred = model.predict(x_test)
```

```
    print(f"\n{name} Metrics:")
```

```
    print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
    print("Precision:", precision_score(y_test, y_pred))
```

```
    print("Recall:", recall_score(y_test, y_pred))
```

```
    print("F1 Score:", f1_score(y_test, y_pred))
```

```
    print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
# Optimization: Tune Random Forest
```

```
tuned_rf = RandomForestClassifier(n_estimators=100, max_depth=5, random_state=42)
```

```
tuned_rf.fit(x_train, y_train)
```

```
y_pred_tuned = tuned_rf.predict(x_test) # Predict on x_test, not x_train
```

```
print("\nTuned Random Forest Metrics:")
```

```
print("Accuracy:", accuracy_score(y_test, y_pred_tuned))
```

```
print("Precision:", precision_score(y_test, y_pred_tuned))
```

```
print("Recall:", recall_score(y_test, y_pred_tuned))
```

```
print("F1 Score:", f1_score(y_test, y_pred_tuned))
```

Logistic Regression Metrics:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 Score: 1.0

Classification Report:

	precision	recall	f1-score	support
1	1.00	1.00	1.00	1
accuracy			1.00	1
macro avg	1.00	1.00	1.00	1
weighted avg	1.00	1.00	1.00	1

Random Forest Metrics:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 Score: 1.0

Classification Report:

	precision	recall	f1-score	support
1	1.00	1.00	1.00	1
accuracy			1.00	1
macro avg	1.00	1.00	1.00	1
weighted avg	1.00	1.00	1.00	1

Tuned Random Forest Metrics:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 Score: 1.0