

```
# Import necessary libraries

import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import classification_report, confusion_matrix

from imblearn.over_sampling import SMOTE

from sklearn.linear_model import LogisticRegression
```

```
# Load the dataset
```

```
# Replace 'path_to_csv' with the actual file path
```

```
data = pd.read_csv('creditcard.csv')
```

```
print(data)
```

```

      Time      V1      V2      V3      V4      V5      V6      V7      ...      V23      V24      V25      V26      V27      V28      Amount      Class
0      0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599  ... -0.110474  0.066928  0.128539 -0.189115  0.133558 -0.021053  149.62      0
1      0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803  ...  0.101288 -0.339846  0.167170  0.125895 -0.008983  0.014724    2.69      0
2      1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461  ...  0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752   378.66      0
3      1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609  ... -0.190321 -1.175575  0.647376 -0.221929  0.062723  0.061458   123.50      0
4      2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941  ... -0.137458  0.141267 -0.206010  0.502292  0.219422  0.215153    69.99      0
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
284802  172786.0 -11.881118  10.071785 -9.834783 -2.066656 -5.364473 -2.606837 -4.918215  ...  1.014480 -0.509348  1.436807  0.250034  0.943651  0.823731    0.77      0
284803  172787.0 -0.732789 -0.055080  2.035030 -0.738589  0.068229  1.058415  0.024330  ...  0.012463 -1.016226 -0.606624 -0.395255  0.068472 -0.053527   24.79      0
284804  172788.0  1.919565 -0.301254 -3.249640 -0.557828  2.630515  3.031260 -0.296827  ... -0.037501  0.640134  0.265745 -0.087371  0.004455 -0.026561    67.88      0
284805  172788.0 -0.240440  0.530483  0.702510  0.689799 -0.377961  0.623708 -0.686180  ... -0.163298  0.123205 -0.569159  0.546668  0.108821  0.104533   10.00      0
284806  172792.0 -0.533413 -0.189733  0.703337 -0.506271 -0.012546 -0.649617  1.577006  ...  0.376777  0.008797 -0.473649 -0.818267 -0.002415  0.013649   217.00      0

[284807 rows x 31 columns]
```

```
# Check for missing values and basic dataset info
```

```
print(f"Dataset shape: {data.shape}")
```

```
[284807 rows x 31 columns]
Dataset shape: (284807, 31)
```

```
print(f"Missing values:\n{data.isnull().sum()}")
```

```
Missing values:
```

```
Time      0
```

```
V1        0
```

```
V2        0
```

```
V3        0
```

```
V4        0
```

```
V5        0
```

```
V6        0
```

```
V7        0
```

```
V8        0
```

```
V9        0
```

```
V10       0
```

```
V11       0
```

```
V12       0
```

```
V13       0
```

```
V14       0
```

```
V15       0
```

```
V16       0
```

```
V17       0
```

```
V18       0
```

```
V19       0
```

```
V20       0
```

```
V21       0
```

```
V22       0
```

```
V23       0
```

```
V24       0
```

```
V25       0
```

```
V26       0
```

```
V27       0
```

```
V28       0
```

```
Amount    0
```

```
Class     0
```

```
dtype: int64
```

```
# Separate features and target
```

```
X = data.drop(columns=['Class']) # Replace 'Class' with your target column name if different
```

```
print(X)
```

```
y = data['Class']
```

```
print(y)
```

```
      Time      V1      V2      V3      V4      V5      V6      V7      ...      V22      V23      V24      V25      V26      V27      V28      Amount
0      0.0 -1.359807 -0.072781 2.536347 1.378155 -0.338321 0.462388 0.239599 ... 0.277838 -0.110474 0.066928 0.128539 -0.189115 0.133558 -0.021053 149.62
1      0.0  1.191857  0.266151 0.166480 0.448154 0.060018 -0.082361 -0.078803 ... -0.638672 0.101288 -0.339846 0.167170 0.125895 -0.008983 0.014724  2.69
2      1.0 -1.358354 -1.340163 1.773209 0.379780 -0.503198 1.800499 0.791461 ... 0.771679 0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752 378.66
3      1.0 -0.966272 -0.185226 1.792993 -0.863291 -0.010309 1.247203 0.237609 ... 0.005274 -0.190321 -1.175575 0.647376 -0.221929 0.062723 0.061458 123.50
4      2.0 -1.158233  0.877737 1.548718 0.403034 -0.407193 0.095921 0.592941 ... 0.798278 -0.137458 0.141267 -0.206010 0.502292 0.219422 0.215153  69.99
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
284802 172786.0 -11.881118 10.071785 -9.834783 -2.066656 -5.364473 -2.606837 -4.918215 ... 0.111864 1.014480 -0.509348 1.436807 0.250034 0.943651 0.823731  0.77
284803 172787.0 -0.732789 -0.055080 2.035030 -0.738589 0.868229 1.058415 0.024330 ... 0.924384 0.012463 -1.016226 -0.606624 -0.395255 0.068472 -0.053527 24.79
284804 172788.0 1.919565 -0.301254 -3.249640 -0.557828 2.630515 3.031260 -0.296827 ... 0.578229 -0.037501 0.640134 0.265745 -0.087371 0.004455 -0.026561 67.88
284805 172788.0 -0.240440 0.530483 0.702510 0.689799 -0.377961 0.623708 -0.686180 ... 0.800049 -0.163298 0.123205 -0.569159 0.546668 0.108821 0.104533 10.00
284806 172792.0 -0.533413 -0.189733 0.703337 -0.506271 -0.012546 -0.649617 1.577006 ... 0.643078 0.376777 0.008797 -0.473649 -0.818267 -0.002415 0.013649 217.00

[284807 rows x 30 columns]
```

```
# Normalize the features
```

```
scaler = StandardScaler()
```

```
print(scaler)
```

```
X_scaled = scaler.fit_transform(X)
```

```
print(X_scaled)
```

```
      Time      V1      V2      V3      V4      V5      V6      V7      ...      V22      V23      V24      V25      V26      V27      V28      Amount
0      0.0 -1.359807 -0.072781 2.536347 1.378155 -0.338321 0.462388 0.239599 ... 0.277838 -0.110474 0.066928 0.128539 -0.189115 0.133558 -0.021053 149.62
1      0.0  1.191857  0.266151 0.166480 0.448154 0.060018 -0.082361 -0.078803 ... -0.638672 0.101288 -0.339846 0.167170 0.125895 -0.008983 0.014724  2.69
2      1.0 -1.358354 -1.340163 1.773209 0.379780 -0.503198 1.800499 0.791461 ... 0.771679 0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752 378.66
3      1.0 -0.966272 -0.185226 1.792993 -0.863291 -0.010309 1.247203 0.237609 ... 0.005274 -0.190321 -1.175575 0.647376 -0.221929 0.062723 0.061458 123.50
4      2.0 -1.158233  0.877737 1.548718 0.403034 -0.407193 0.095921 0.592941 ... 0.798278 -0.137458 0.141267 -0.206010 0.502292 0.219422 0.215153  69.99
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
284802 172786.0 -11.881118 10.071785 -9.834783 -2.066656 -5.364473 -2.606837 -4.918215 ... 0.111864 1.014480 -0.509348 1.436807 0.250034 0.943651 0.823731  0.77
284803 172787.0 -0.732789 -0.055080 2.035030 -0.738589 0.868229 1.058415 0.024330 ... 0.924384 0.012463 -1.016226 -0.606624 -0.395255 0.068472 -0.053527 24.79
284804 172788.0 1.919565 -0.301254 -3.249640 -0.557828 2.630515 3.031260 -0.296827 ... 0.578229 -0.037501 0.640134 0.265745 -0.087371 0.004455 -0.026561 67.88
284805 172788.0 -0.240440 0.530483 0.702510 0.689799 -0.377961 0.623708 -0.686180 ... 0.800049 -0.163298 0.123205 -0.569159 0.546668 0.108821 0.104533 10.00
284806 172792.0 -0.533413 -0.189733 0.703337 -0.506271 -0.012546 -0.649617 1.577006 ... 0.643078 0.376777 0.008797 -0.473649 -0.818267 -0.002415 0.013649 217.00

[284807 rows x 30 columns]
```

```
# Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42, stratify=y)
```

```
print(X_train,)
```

```
StandardScaler()  
[[ -1.99658302 -0.69424232 -0.04407492 ...  0.33089162 -0.06378115  
   0.24496426]  
 [ -1.99658302  0.60849633  0.16117592 ... -0.02225568  0.04460752  
  -0.34247454]  
 [ -1.99656197 -0.69350046 -0.81157783 ... -0.13713686 -0.18102083  
   1.16068593]  
 ...  
 [  1.6419735   0.98002374 -0.18243372 ...  0.01103672 -0.0804672  
  -0.0818393 ]  
 [  1.6419735  -0.12275539  0.32125034 ...  0.26960398  0.31668678  
  -0.31324853]  
 [  1.64205773 -0.27233093 -0.11489898 ... -0.00598394  0.04134999  
   0.51435531]  
 [[  1.41309493  0.99390119 -0.45571564 ...  0.19235561 -0.09769602  
  -0.3239634 ]  
 [  0.62464415  1.03903462 -0.02960101 ...  0.00624303 -0.20904363  
  -0.34127512]  
 [ -1.12918609 -0.50641931  0.36528224 ...  0.09426583  0.56149664  
   0.3468355 ]  
 ...  
 [ -1.25321726  0.54616878  0.04366533 ...  0.07683742  0.07539328  
  -0.17163632]  
 [ -1.48249593  0.65373478  0.18202927 ... -0.14824131 -0.0178351  
  -0.3496711 ]  
 [ -1.3891251  -0.30536718  0.46935012 ...  0.9379349   0.62519565  
  -0.32528277]]
```

```
# Handle class imbalance using SMOTE
```

```
smote = SMOTE(random_state=42)
```

```
print(smote)
```

```
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

```
print(X_train_resampled,y_train_resampled)
```

```

1      0
2      0
3      0
4      0
..
454897  1
454898  1
454899  1
454900  1
454901  1
Name: Class, Length: 454902, dtype: int64

```

```

# Choose a classification algorithm
# Use Logistic Regression for faster training
# You can switch to RandomForestClassifier if needed
classifier = LogisticRegression(max_iter=500, random_state=42)
print(classifier)

```

```

LogisticRegression(max_iter=500, random_state=42)

```

```

# Train the model
print("Training the model...")
print(classifier.fit(X_train_resampled, y_train_resampled))

```

```

Training the model...
LogisticRegression(max_iter=500, random_state=42)

```

```

# Make predictions on the test set
y_pred = classifier.predict(X_test)
print(y_pred)

```

```

[0 0 0 ... 0 0 0]

```

```
# Evaluate the model's performance

print("\nConfusion Matrix:")

print(confusion_matrix(y_test, y_pred))
```

```
Confusion Matrix:
[[55399  1465]
 [      8    90]]
```

```
print("\nClassification Report:")

print(classification_report(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0           1.00      0.97      0.99     56864
     1           0.06      0.92      0.11         98

 accuracy          0.97          0.97     56962
 macro avg         0.53          0.95      0.55     56962
 weighted avg      1.00          0.97      0.99     56962
```

Calculate precision, recall, and F1-score

```
from sklearn.metrics import precision_score, recall_score, f1_score

precision = precision_score(y_test, y_pred)

recall = recall_score(y_test, y_pred)

f1 = f1_score(y_test, y_pred)

print(f"\nPrecision: {precision:.2f}")

print(f"Recall: {recall:.2f}")

print(f"F1-Score: {f1:.2f}")
```

```
Precision: 0.06
Recall: 0.92
F1-Score: 0.11
```

```
import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.metrics import roc_curve, auc, precision_recall_curve
```

```
# Plot 1: Class Distribution5

plt.figure(figsize=(6, 4))


sns.countplot(x=y)

plt.title("Class Distribution")

plt.xlabel("Class (0: Genuine, 1: Fraudulent)")

plt.ylabel("Count")

plt.show()
```



```
# Plot 2: Correlation Heatmap
```


```
plt.figure(figsize=(12, 10))
```

```
correlation_matrix = data.corr()
```

```
sns.heatmap(correlation_matrix, cmap='coolwarm', annot=False)
```

```
plt.title("Feature Correlation Heatmap")
```

```
plt.show()
```



```
# Plot 3: Feature Distribution (After Scaling)
```

```
plt.figure(figsize=(14, 8))
```

```
for i, col in enumerate(X.columns[:10], 1): # Adjust for number of features to visualize
```

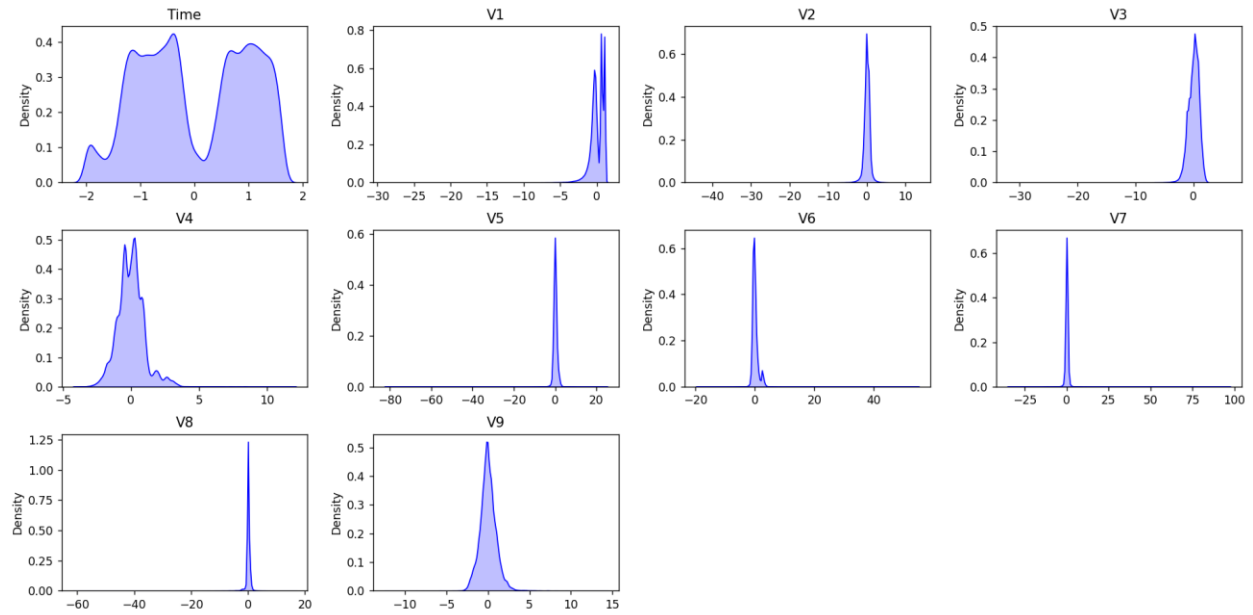
```
    plt.subplot(3, 4, i)
```

```
    sns.kdeplot(X_scaled[:, i - 1], color='blue', fill=True)
```

```
    plt.title(col)
```

```
plt.tight_layout()
```

```
plt.show()
```

Plot 4: ROC Curve

```
y_proba = classifier.predict_proba(X_test)[:, 1] # Probability scores for the positive class
```

```
print(y_proba)
```

```
fpr, tpr, _ = roc_curve(y_test, y_proba)
```


```
print(fpr, tpr)
```

```
roc_auc = auc(fpr, tpr)
```

```
print(roc_auc)
```



```
plt.figure(figsize=(8, 6))  
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC Curve (AUC = {roc_auc:.2f})')  
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')  
plt.xlabel("False Positive Rate")  
plt.ylabel("True Positive Rate")  
plt.title("Receiver Operating Characteristic (ROC) Curve")  
plt.legend(loc="lower right")  
plt.show()
```



Plot 5: Precision-Recall Curve

```
precision, recall, _ = precision_recall_curve(y_test, y_proba)

plt.figure(figsize=(8, 6))

plt.plot(recall, precision, color='blue', lw=2)

plt.xlabel("Recall")

plt.ylabel("Precision")

plt.title("Precision-Recall Curve")

plt.show()
```

